



Comet

Moving towards a real-time Web

Simon Willison
Highland Fling
3rd April 2008



Comet

~~Moving towards a real-time Web~~
Real developers fight browsers with their bare hands!

Simon Willison
Highland Fling
3rd April 2008

Comet: any technique
that allows a web server
to push events “live” to
connected browsers

A new name for a very old technique

- Also known as...
 - HTTP push
 - Server push
 - Reverse Ajax
 - Pushlets

[Home](#) | [Archives](#) | [Contact](#)

Continuing Intermittent Incoherency

Alex Russell being boring, when comprehensibleFri 3
Mar
2006

Comet: Low Latency Data for the Browser

Posted by alex under [programming](#), [javascript](#), [webdev](#)

An old web technology is slowly being resurrected from the depths of history. Browser features that have gone untouched for years are once again being employed to bring better responsiveness to UIs. Servers are learning to cope with a new way of doing things. And I'm not talking about Ajax.

New services like [Jot Live](#) and [Meebo](#) are built with a style of data transmission that is neither traditional nor Ajax. Their brand of low-latency data transfer to the browser is unique, and it is becoming ever-more common. Lacking a better term, I've taken to calling this style of event-driven, server-push data streaming "Comet". It doesn't stand for anything, and I'm not sure that it should. There is much confusion about how these techniques work, and so using pre-existing definitions and names is as likely to get as much wrong as it would get right.

Archived Entry

Post Date :Friday, Mar 3rd, 2006 at
10:37 pm**Category :**[programming](#) and
[javascript](#) and [webdev](#)**Do More :**You can [leave a response](#),
or [trackback](#) from your own
site.

Defining Comet
Our next round of useful, attainable web standards examples focus on a new technology, a list of the problems being solved, and properties which distinguish it from other techniques. As with Ajax, these aren't hard to find.



Deliberately named after a kitchen cleaner

What can you do with it?

Google Mail - Inbox (6655) - swillison@gmail.com

<https://mail.google.com/mail/?shva=1#inbox>

Mail Calendar Documents Photos Groups Web more ▾ [Settings](#) | [Older version](#) | [Help](#) | [Sign out](#)

Zimbra Chris Smith

Compose Mail

Inbox (6655)

Starred

Chats

Sent Mail

Drafts (9)

All Mail

Spam (1414)

Trash

Contacts

Chat

Search, add

Simon W

Set status

Derek H

Mark Pilg

Nathanie

Ryan Kin

@home

damian

Give it up

Dori Smi

too busy

Natalie C

Away

Scott Joh

Code ! C

Stuart La

doing wo

Waiting for m

February

S	M	T	W
27	28	29	30
3	4	5	6
10	11	12	13

New Buddy

Filter

Buddy List

Zimbra As

Buddies

AmyA

GregA

sdtietze

Search

Email ▾

Search

Save

Advanced

Search the Web...

meebo.com

http://www1.meebo.com/

Unsaved spreadsheet - Google Docs

http://spreadsheets.google.com/ccc?key=pAj7BSLZNsDBjyawHiFckDg&hl=en

swillison@gmail.com | Docs Home | Help | Sign Out

Save

Save & close

Preview

Print

Discuss

Share

Publish

File ▾

Edit

Sort

Formulas

Revisions

Format ▾

B

I

U

Abc

F

T

T

Align ▾

Insert ▾

Delete ▾

Wrap Text

Merge across

A B C D E F G H I J K L M N O

1

2

3

4

5

6

7

8

9

10

11

12

13

14

15

16

17

18

19

20

21

22

23

24

25

26

Add Sheet

Sheet1 ▾

The image captures a multitasking session on a Mac OS X desktop. In the foreground, a Google Docs spreadsheet is open, showing a blank 16x16 grid. Behind it, a Zimbra interface displays a list of contacts and buddies. To the right, a Meebo browser window shows a single search result for 'meebo.com'. In the top left, a Gmail inbox window lists several unread messages. The desktop background is black, and the overall layout is characteristic of a mid-2000s Mac OS X desktop environment.

example spreadsheet – Google Docs

http://spreadsheets.google.com/ccc?key=p0GUf1c5U9mzr3Ey5pnQvYA&hl=en_GB

Google Docs BETA

example spreadsheet

File Edit Sort F

Format

	A	B
1		
2		
3		
4		
5		
6		
7		
8		
9		
10		
11		
12		

example spreadsheet – Google D

http://spreadsheets.google.com/ccc?key=p0GUf1c5U9mzr3Ey5pnQvYA

Google Docs BETA

example spreadsheet Autosaved at 4 Nov 2007 09:06:13 GMT

File Edit Sort Formulas Revisions

Format

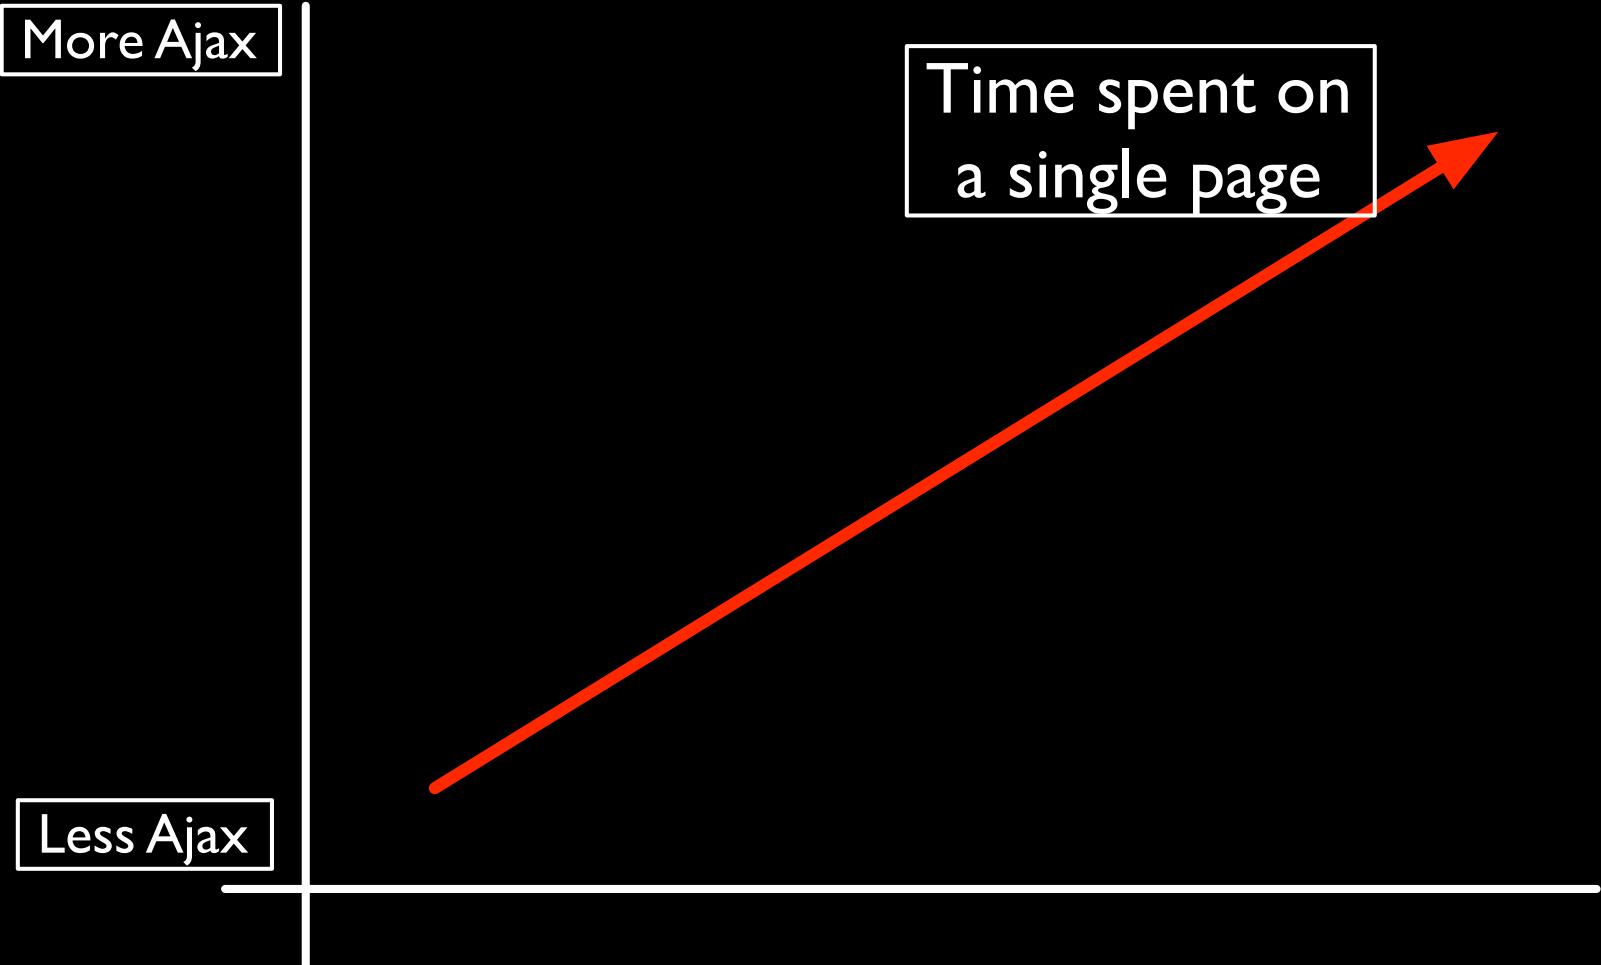
B I U Abc f tT T Align

	A	B	C	D	E	F	G
1							
2							
3							
4							
5							
6							
7							
8							
9							
10							
11							
12							

More Ajax

Less Ajax

Time spent on
a single page

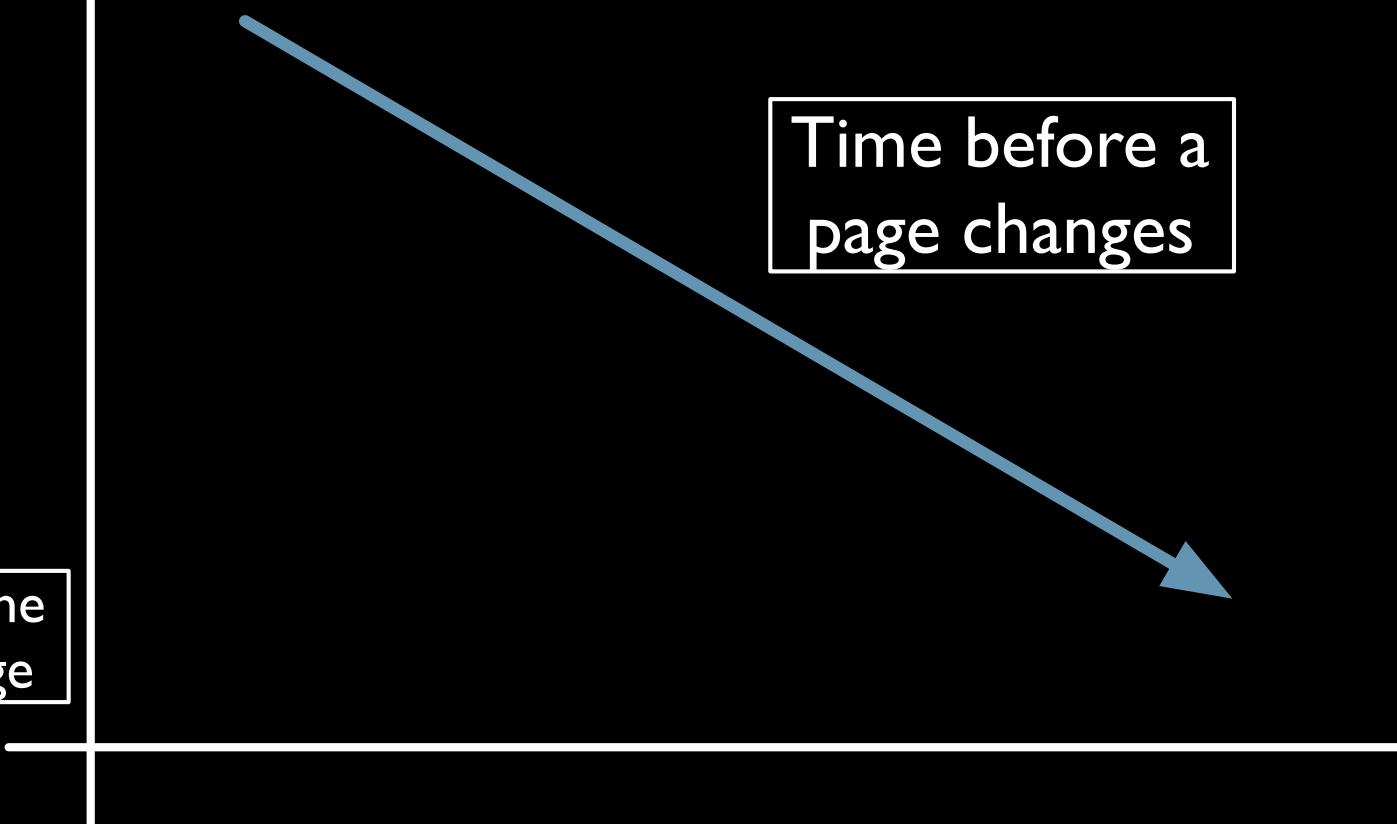


(Slides courtesy of Joe Walker)

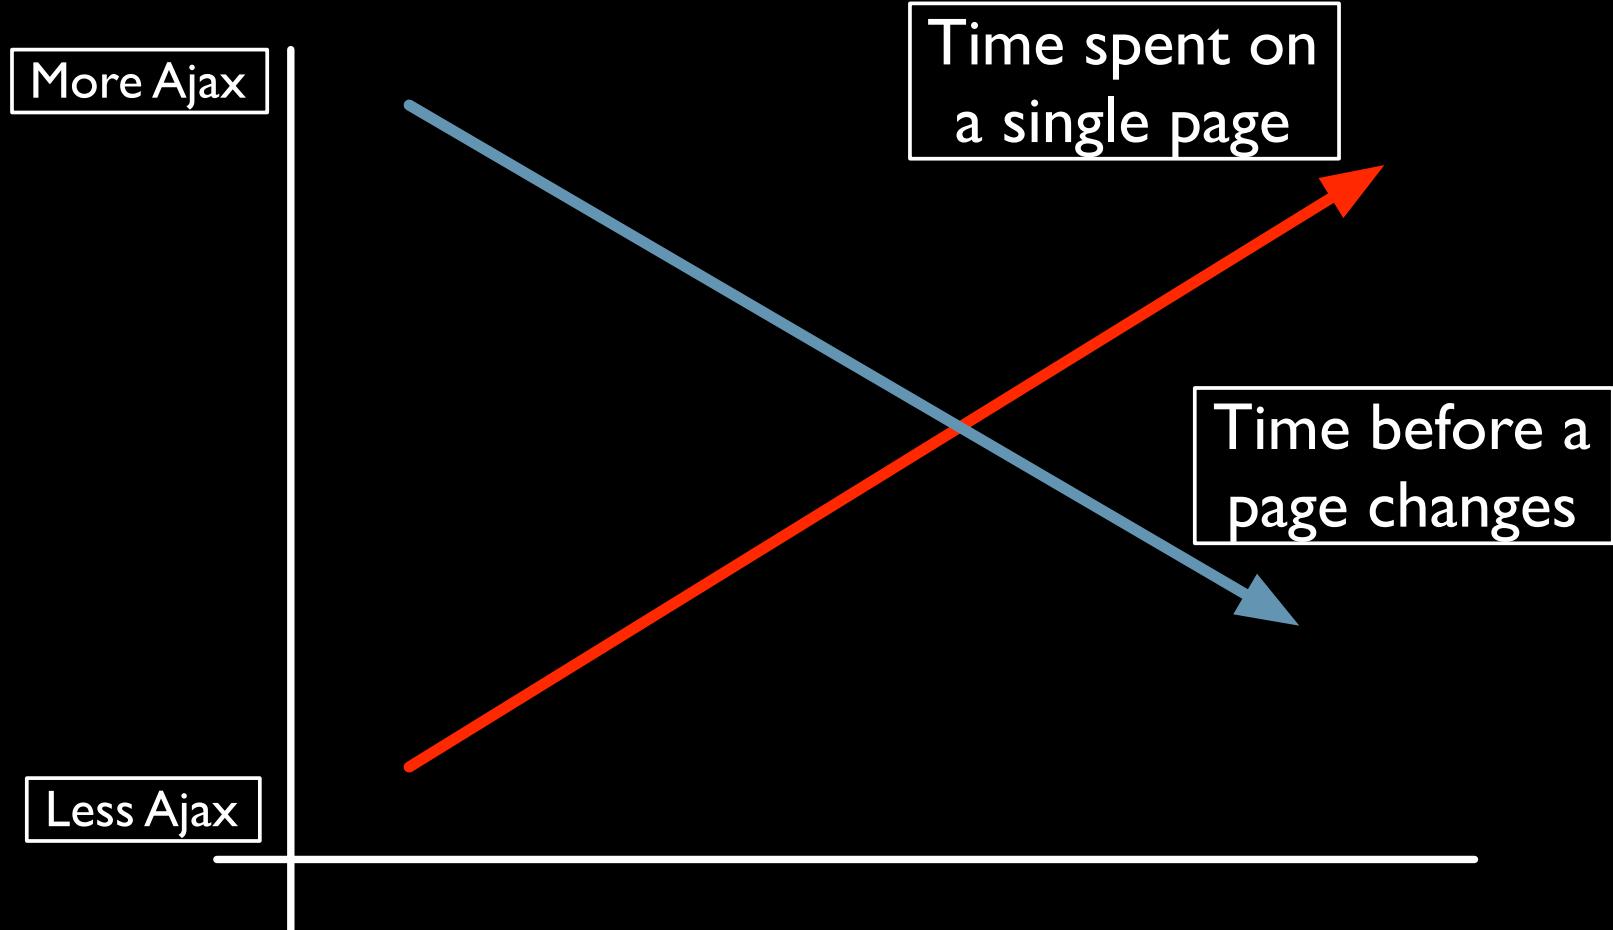
Long time
to change

Short time
to change

Time before a
page changes



(Slides courtesy of Joe Walker)



(Slides courtesy of Joe Walker)

Early Comet



Back



Forward



Home



Reload



Images



Open



Print



Find



Stop

Location:

about:mozilla



What's New!

What's Cool!

Handbook

Net Search

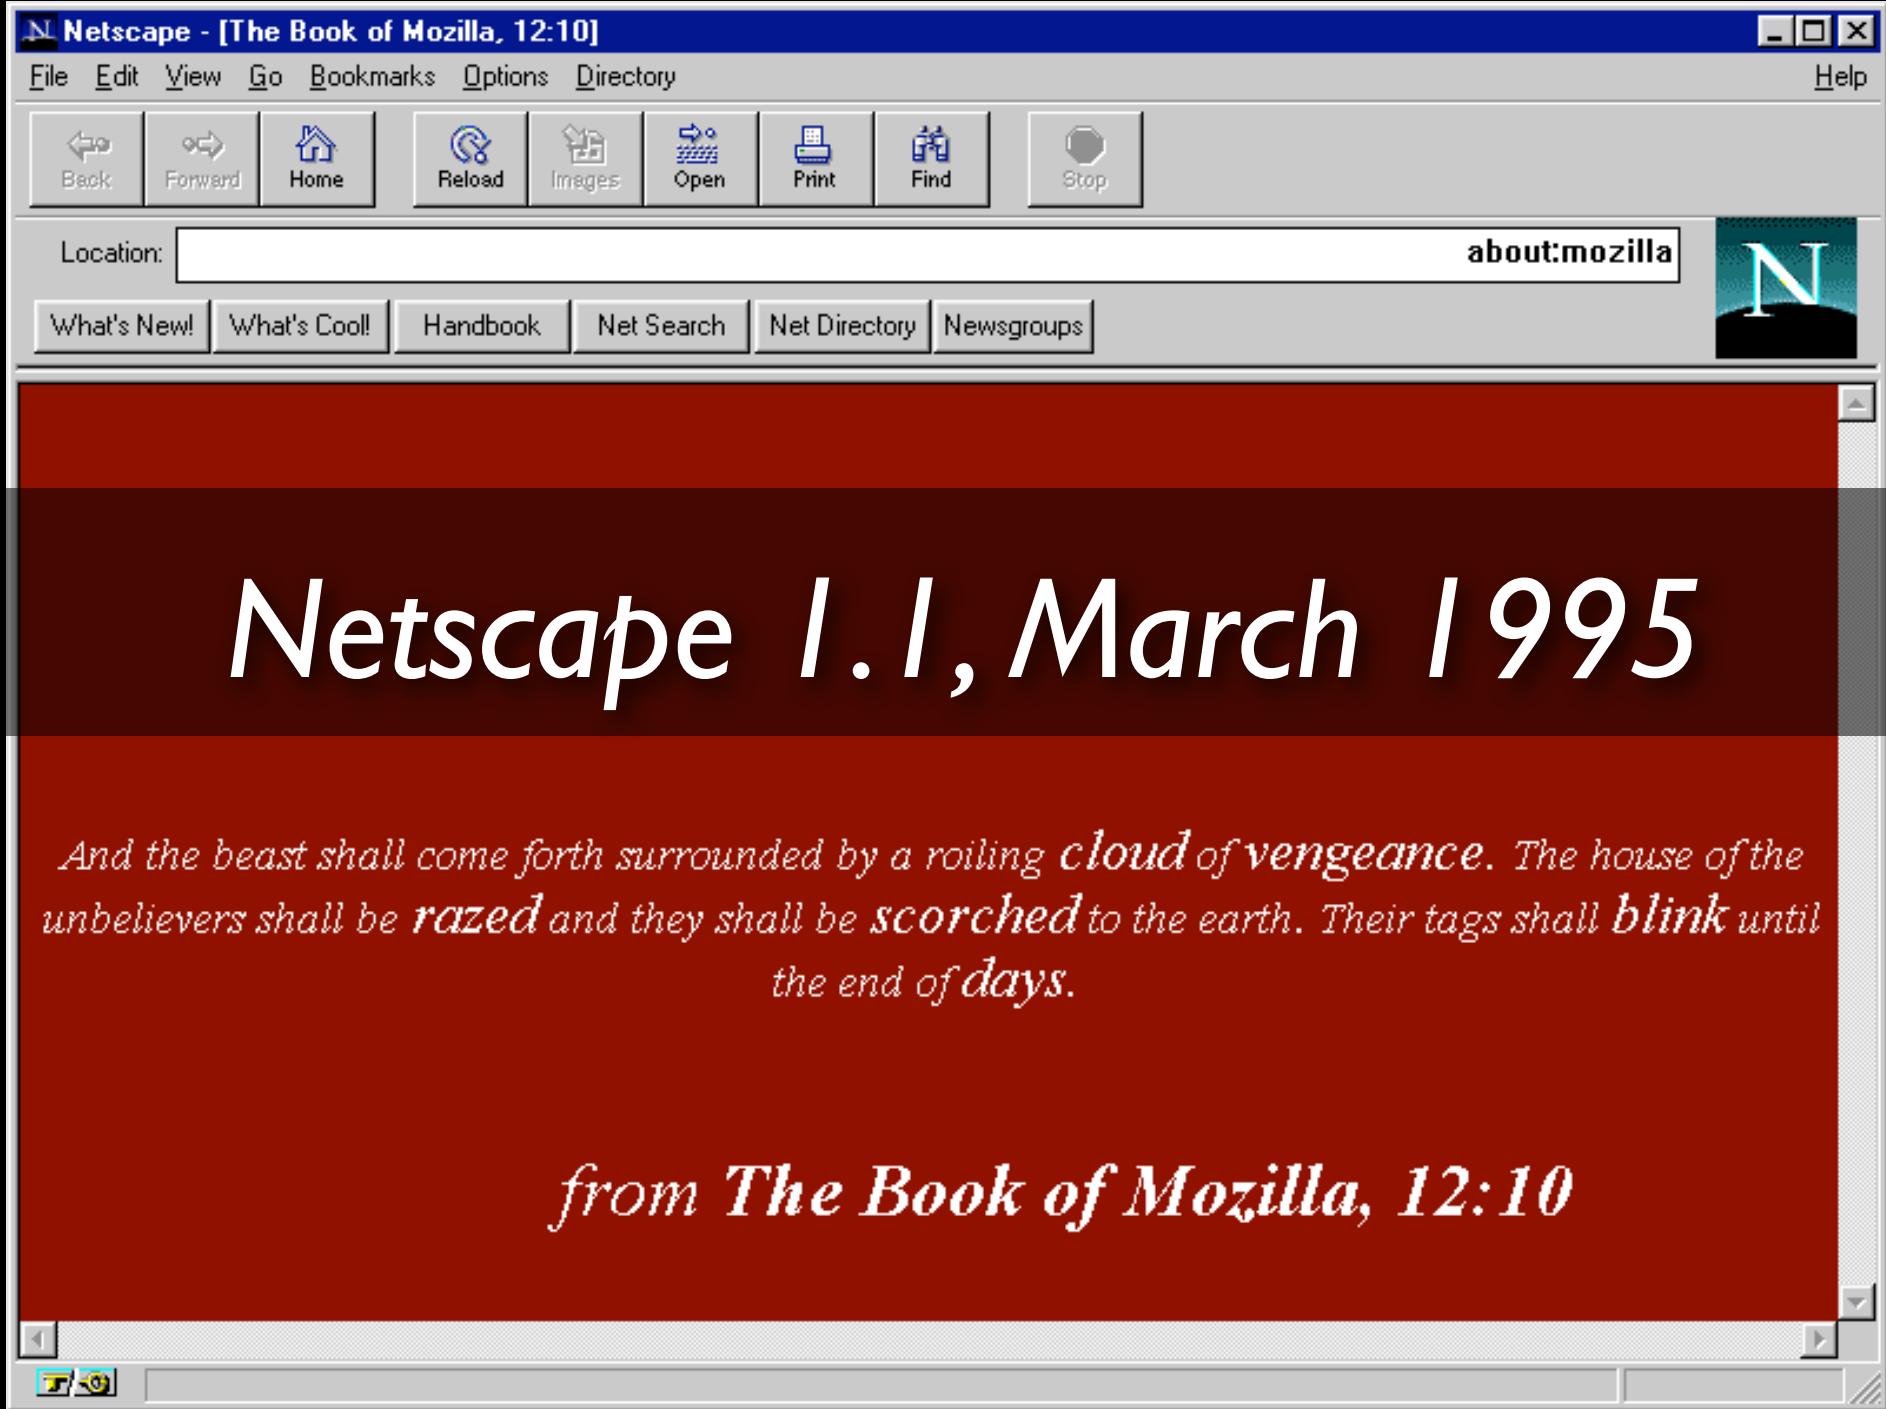
Net Directory

Newsgroups

And the beast shall come forth surrounded by a roiling cloud of vengeance. The house of the unbelievers shall be razed and they shall be scorched to the earth. Their tags shall blink until the end of days.

from The Book of Mozilla, 12:10





Netscape 1.1 new features

- Tables
- Background images
- “Dynamic Documents”
 - Client Pull
 - **Server Push**
- Predates JavaScript, Internet Explorer and Flash



AN EXPLORATION OF DYNAMIC DOCUMENTS

INTRODUCTION

This document explores the "server push" and "client pull" dynamic document capabilities of Netscape Navigator 1.1 (Windows, Mac, and Unix versions). Please send comments to pushpull@netscape.com. Also, if you are using either server push or client pull in a real-world application, please drop us a note at pushpull@netscape.com and let us know -- thanks!

Examples are given at the end, along with an important note on implementing server push CGI programs as shell scripts.

THE GREAT IDEA

The general idea is that browsers have always been driven by user input. You click on a link or an icon or an image and some data comes to you. As soon as people saw they could do that, they wanted to give a server the ability to push new data down to the browser. (An obvious example is a stock trader who wants to see new quote data every 5 minutes.) Up until now, that hasn't been possible.

Netscape Navigator 1.1 gives content creators and server administrators two new open standards-based mechanisms for making this work. The mechanisms are similar in nature and effect, but complementary. They are:

1. **Server push** -- the server sends down a chunk of data; the browser displays the data but leaves the connection open; whenever the server wants it sends more data and the browser displays it, leaving the connection open; at some later time the server sends down yet more data and the browser displays it; etc.
2. **Client pull** -- the server sends down a chunk of data, including a directive (in the HTTP response or the document header) that says "reload this data in 5 seconds", or "go load this other URL in 10 seconds". After the specified amount of time has elapsed, the client does what it was told -- either reloading the current data or getting new data.

In server push, a HTTP connection is held open for an indefinite period of time (until the server knows it is done sending data to the client and sends a terminator, or until the client disrupts the connection). In client pull, HTTP connections are never held open; rather, the client is told when to open a new connection, and what data to fetch when it does.

Server Push, Client Pull

In server push, the magic is accomplished by using a variant of the MIME message format "multipart/mixed", which lets a single message (or HTTP response) contain many data items. In client pull, the magic is accomplished by an HTTP response header (or equivalent HTML tag) that tells the client what to do after some specified time delay.

“Client Pull”

```
<META HTTP-EQUIV="Refresh" CONTENT=1>
```

“Server Push”

Content-type: **multipart/x-mixed-replace;boundary=XXooXX**

--XXooXX

Content-type: text/plain

Data for the first object.

--XXooXX

Content-type: text/plain

Data for the second and last object.

--XXooXX--



• RESERVATIONS

• CONTACT US

• WEBCAM

• HOME

• PACKAGES

• DINING

• ROOMS & SUITES

• VIEW ROOMS

• ATTRACTIONS

Web Cam

Adobe Resort
PO Box 219,
1555 Hwy 101,
Yachats, OR 97498.

Reservations:
1-800-522-3623
Or
(541) 547-3141
Fax: (541) 547-4234
or [email us](#).

[DIRECTIONS](#)



Enter Address Here



TELL A FRIEND

RESET

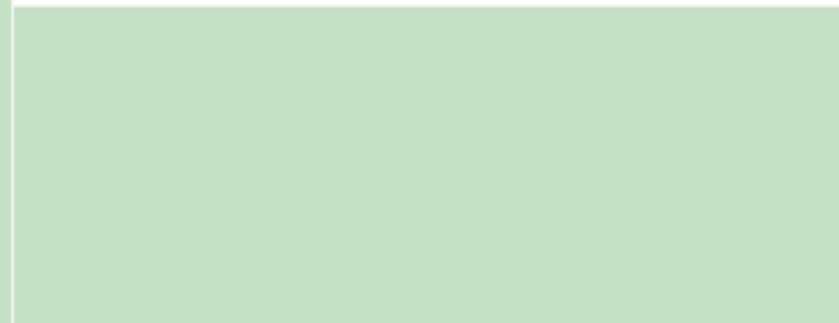
Copyright 2004 by Adobe Resort



Chat Using Dynamic Animated Images

This page demonstrates how text can be typeset and pushed to a web browser in real time. The chat box below is an interactive, dynamically generated, animated GIF image. Each time a line of text is to be displayed, the server renders and sends the next frame of the animation.

```
12:13: teets - dada
12:13: teets - ddaaad
Ping joined at 20:35 on 2007-12-03.
20:35: Ping - Hi Simon
ramana joined at 00:57 on 2007-12-04.
00:57: ramana - hi
Caca joined at 04:04 on 2007-12-04.
04:04: Caca - Salut
04:07: Caca - ca va?
04:07: Caca - Fuck you!
dave joined at 14:37 on 2007-12-04.
14:37: dave - hello
hi joined at 14:41 on 2007-12-04.
hi left at 14:41 on 2007-12-04.
Simon joined at 22:00 on 2007-12-04.
```



Simon:

Say it

This technique makes it possible to provide a real-time chat service with no software for users to download, no firewall or proxy problems, no dynamic HTML, no Java or Javascript required, no Flash, and no automatically reloading frames. This would have worked long before Netscape, Sun, and Microsoft introduced their buggy and bloated implementations of these features.

This little look-alike window in the niche card was built on Saturday 25 September 1999. The font is the product of some early word-processor editing.

Javascript is used for displaying the text in a window-like fashion, except when you close this window, but the Javascript isn't necessary in order to make text typed by others appear in real-time. Sorry I didn't get around to implementing word-wrapping.

<http://zesty.ca/chat/> (1999)

Modern Comet

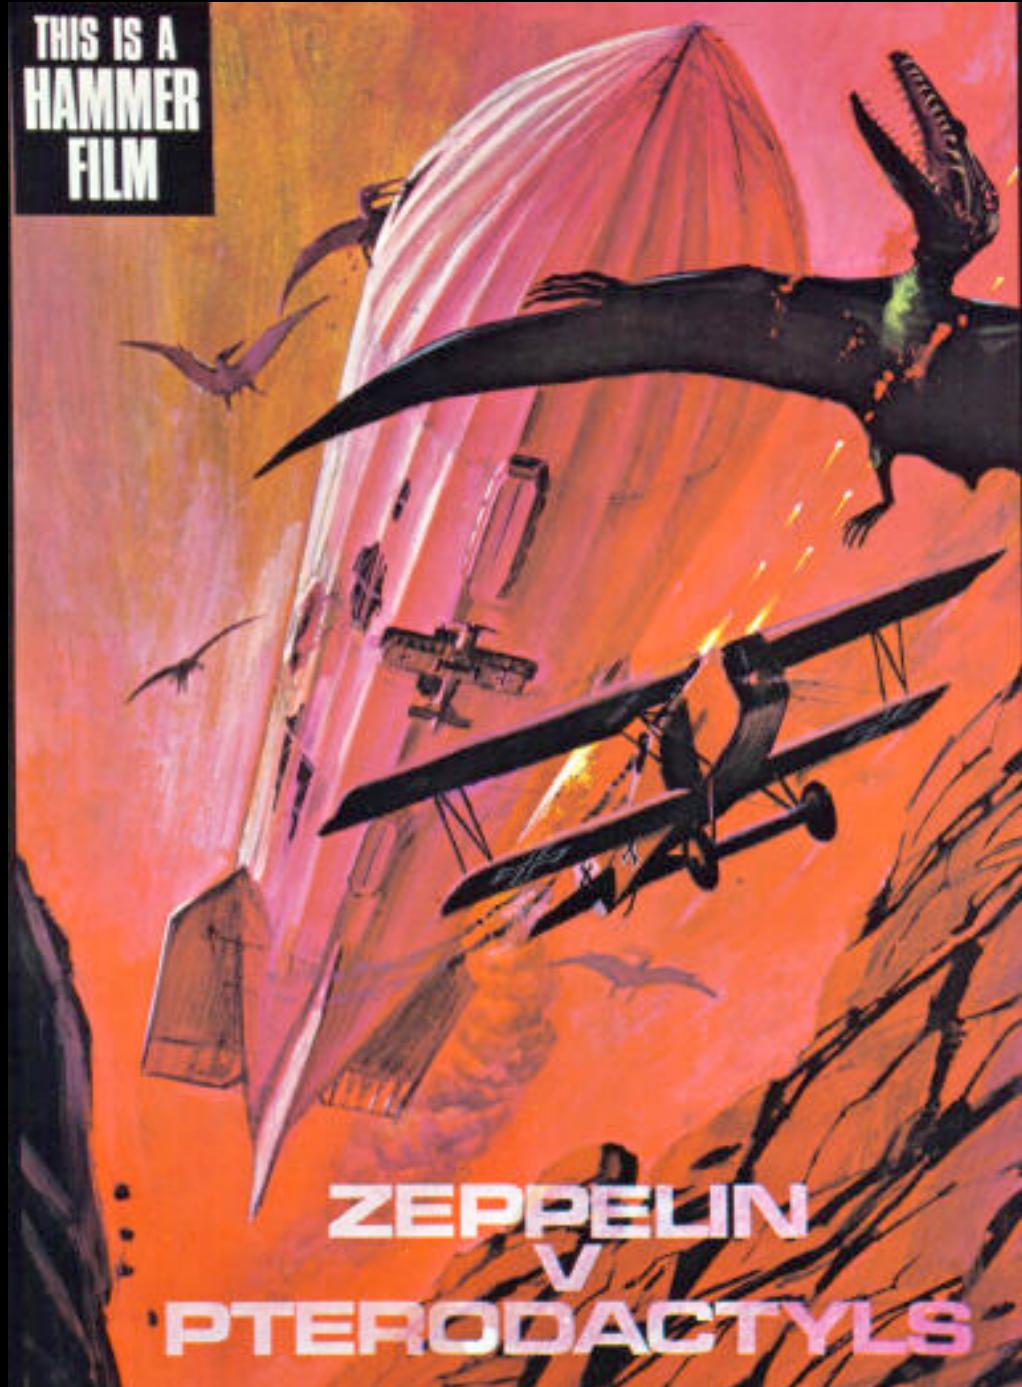
```
var xhr = new XMLHttpRequest();
xhr.open('GET', '/comet', true);

xhr.onreadystatechange = function() {
    if (readyState == 3) {
        // Process xhr.responseText
    }
};
```

If only it were
that simple...

An Epic Saga

Of geek v.s. browser



<http://airminded.org/2007/05/25/the-movie-that-time-forgot/>

Problem #1



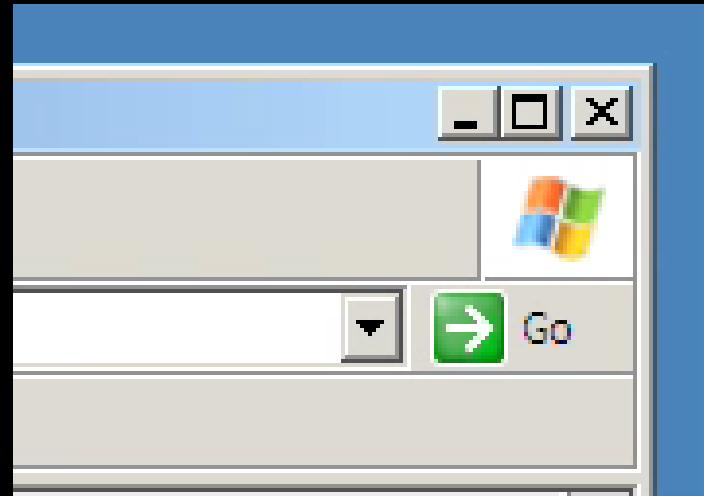
`xhr.responseText` isn't available until
the page has finished loading

A hidden iframe

(The forever frame technique)

- Add a hidden iframe pointing at a streaming page
- Send down chunks of <script> blocks, which the browser will execute as part of progressive rendering
- Add 1KB of junk at the start to trigger progressive rendering in IE

Problem #2



The throbber!

Problem #3

The click!

Solved for GChat

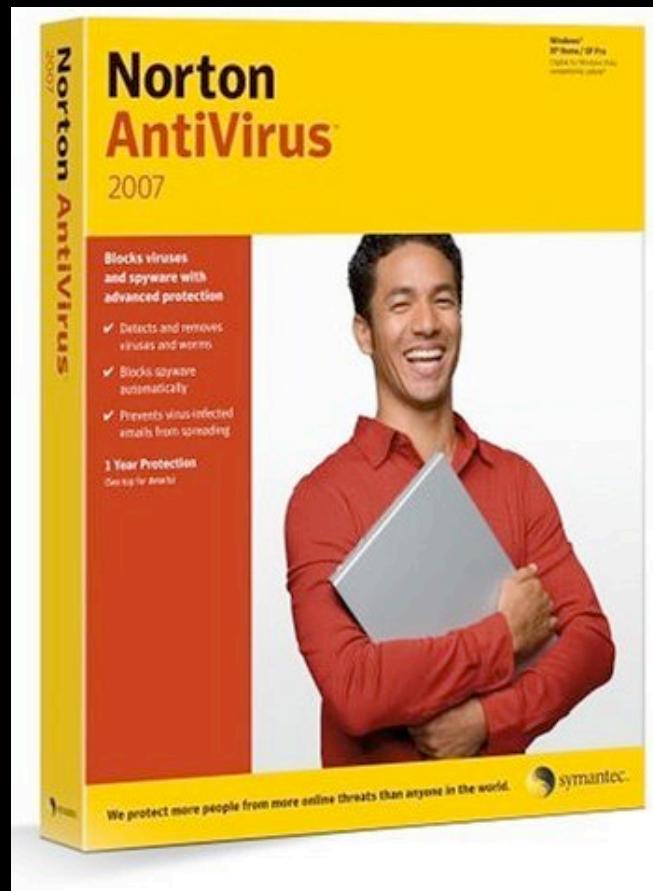
```
var htmlfile = new ActiveXObject('htmlfile');
```

- An undocumented ActiveX extension
 - Create an "htmlfile" instance
 - Create an iframe inside it
 - Set up some cross-frame scripting hooks
 - Run comet using the iframe
- No clicks! No throbber!

Confuse IE with enough nested iframes and it forgets to click



Problem #4



Proxies!

Long polling

- Make a request (iframe, XHR or even a <script> tag)
- Server pretends to be processing it really slowly...
- An event occurs; the server finally returns the response
- The browser starts a new request straight away

Problem #5

The two connection limit

- `mysite.com` - regular site
- `comet.mysite.com` - comet server
- Cross-domain communication is easy - just add another iframe and set `document.domain` to enable cross-frame scripting

- `mysite.com` - regular site
- `comet.mysite.com` - comet server
- Cross-domain communication is easy - just add another iframe and set `document.domain` to enable cross-frame scripting
- (Then solve the resulting browser history problems with a few more iframes)

And if all else fails...



Fall back to polling

THIS IS A
HAMMER
FILM

Pterodactyls win!

ZEPPELIN
V
PTEROVACTYLS

Key techniques

- Streaming (ideal)
 - XMLHttpRequest
 - Forever frame (for IE)
- Long polling (works through proxies)
- Polling (if all else fails)

Future Comet

HTML 5 event-source

<http://www.whatwg.org/specs/web-apps/current-work/#event-source>

```
<event-source src="/comet">
```

Content-Type: **application/x-dom-event-stream**

Event: server-time
data: [time on the server]



Support added in Opera 9.5

Google Gears

http://gears.google.com/ Google

Google Gears™ BETA Improving Your Web Browser

Google Gears is an open source project that enables more powerful web applications, by adding new features to your web browser:

-  Let web applications interact naturally with your desktop
-  Store data locally in a fully-searchable database
-  Run JavaScript in the background to improve performance

Install Google Gears

Mac OS X

System requirements

- Mac OS X 10.2+
- Firefox 1.5+

Google Gears is available for [Windows](#), [Windows Mobile](#), [Mac](#), and [Linux](#)

[Developer site](#)
[Frequently asked questions](#)
[More help](#)

Google Gears 0.2

©2008 Google - [Gears Home](#) - [Privacy Policy](#)

Done

“Wow, client-side
Comet sucks”...

that’s not the half of it



Scaling the server

<http://flickr.com/photos/adrianblack/254968866/> craiglblack

- Comet requires potentially tens of thousands of simultaneous HTTP connections
- Apache and other mainstream servers are designed to handle a single request as quickly as possible
 - Requests are generally served by a worker process or thread
 - This absolutely will not scale to Comet

Event-based IO

- Alternative architecture designed to handle thousands of simultaneous connections
 - Twisted (Python)
 - Jetty (Java)
 - Perlbal
 - Orbited
 - Meteor ...

HAROLD D:

Bayeux

Bayeaux

- Bayeaux is a *protocol* for Comet
- Any Bayeaux client can talk to any Bayeaux server
- The protocol is based on clients publishing and subscribing to channels
- Data is encoded using JSON
- Clients can create a permanent connection using a handshake, or send simple one-off messages

Dumb Comet servers

- The principle advantage of Bayeaux is that your Comet server can be **stupid** - it doesn't need any application logic, it just needs to route messages around
- It's a black box which you simply drop in to your architecture
- This means your application logic can stay on your favourite platform

Cometd

- Cometd (Comet Daemon) is an umbrella project for a number of Bayeaux implementations
 - cometd-python (in Twisted)
 - cometd-perl (also a Perlbal plugin)
 - cometd-java (on Jetty)
 - dojox.cometd (JavaScript client)
- <http://cometd.com/>

So despite all of that...
Comet applications
are *easy* to build

How to build a Comet application (in 5 minutes)

Set up Jetty

- Download and unzip Jetty-6.1 from www.mortbay.org
- Install Maven from <http://maven.apache.org/>
- cd jetty-6.1.6/contrib/cometd/demo/
- mvn jetty:run
- Browse to <http://localhost:8080/>
- Mess around with the JS in src/main/webapp/examples/

dojox.cometd

```
dojo.addOnLoad(function() {
    dojox.cometd.init("http://127.0.0.1:8080/cometd");
    dojox.cometd.subscribe("/mychannel", function(comet) {
        alert(comet.data); // a JSON object
    });
});
```

dojox.cometd

```
dojo.addOnLoad(function() {
    dojox.cometd.init("http://127.0.0.1:8080/cometd");
    dojox.cometd.subscribe("/mychannel", function(comet) {
        alert(comet.data); // a JSON object
    });
}

dojo.byId('send').onclick = function() {
    dojox.cometd.publish("/mychannel", {
        'msg': "A message sent by Comet"
    });
    return false;
};
});
```

A slideshow in 5 lines of code

```
<script type="text/javascript">
dojo.require("dojox.cometd");
jQuery(function($) {
    dojox.cometd.init("http://127.0.0.1:8080/cometd");
    dojox.cometd.subscribe("/slideshow/change", function(comet) {
        $('#currentSlide').attr('src', comet.data.src);
    });
});
</script>
```

Conclusions

- Comet requires hacks, but they're reasonably well understood
- Bayeaux and Cometd abstract away the nastiness and make Comet accessible to sane developers
- If we survived CSS, Comet should be a cinch

More crazy hacks...

- <http://meteorserver.org/browser-techniques/>
- <http://cometdaily.com/2007/10/25/http-streaming-and-internet-explorer/>
- <http://cometdaily.com/2007/11/18/ie-activexhtmlfile-transport-part-ii/>
- <http://orbited.org/svn/orbit/trunk/daemon/orbited/static/orbited.js>
- <http://simonwillison.net/tags/comet/>

Thank you